Course: "Introduction to Computational Thinking and Data Science"

Yolanda Gil
gil@isi.edu

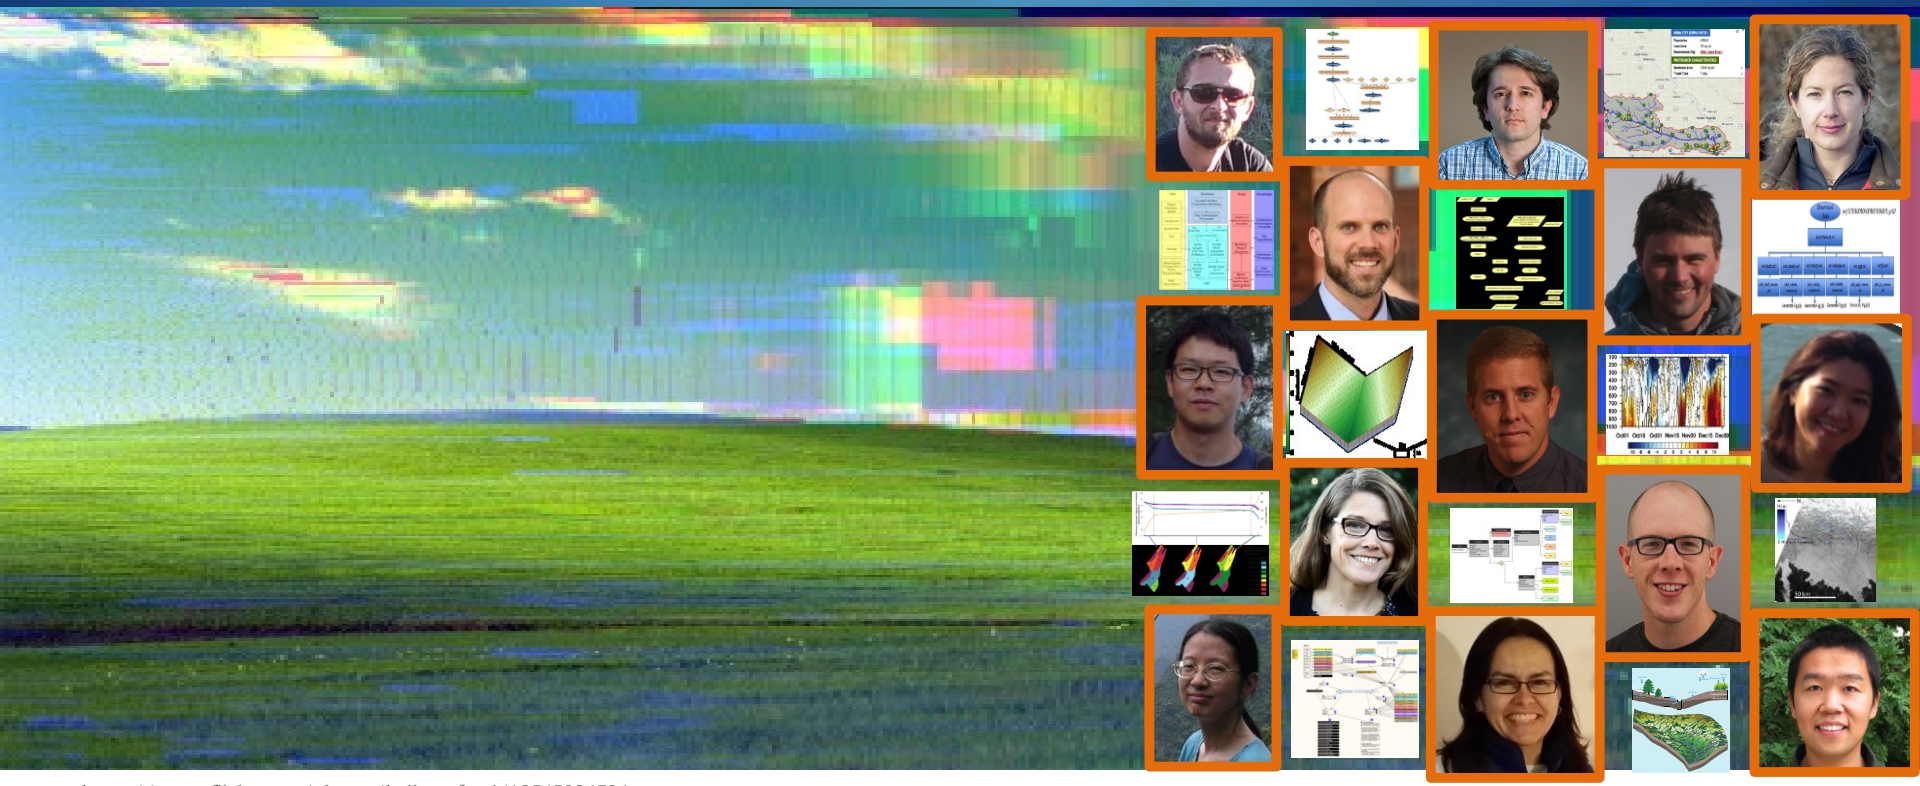

https://www.flickr.com/photos/hellocatfood/10545396704

Goals of the Course

Course is **designed for students with no programming background who want to have literacy in computing and data science** to better approach data-rich problems

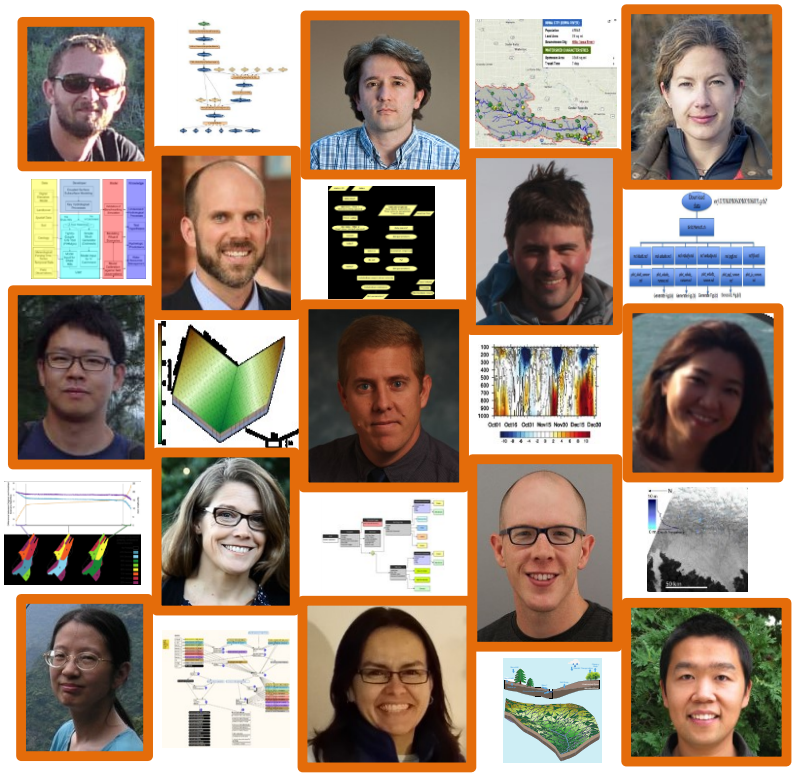

Barriers of Non-Programmers to Data Science

Describe problem

Provide data

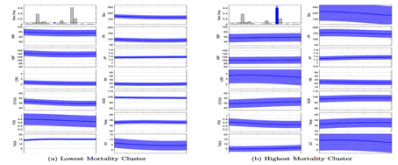

Show results

Point out issues

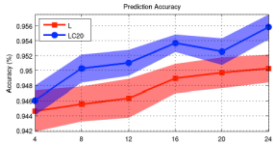

Show more results

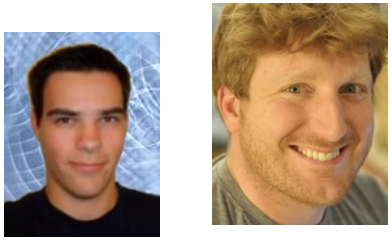

CODE

Distinct Expertise in Data Science

Domain knowledge

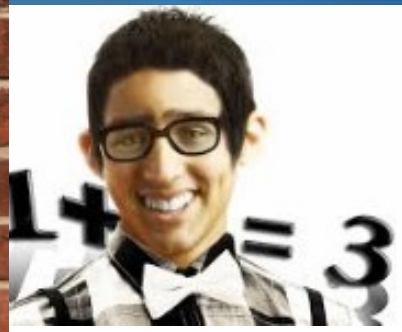

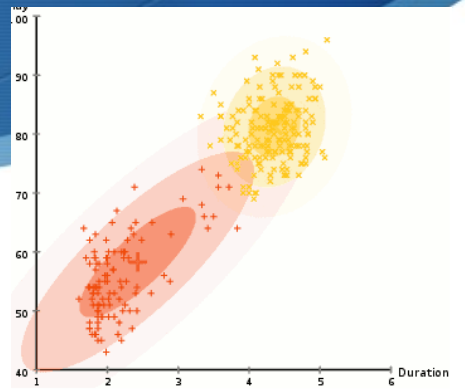

Statistics, data mining

Semantics and distributed data

Large-Scale Data Processing

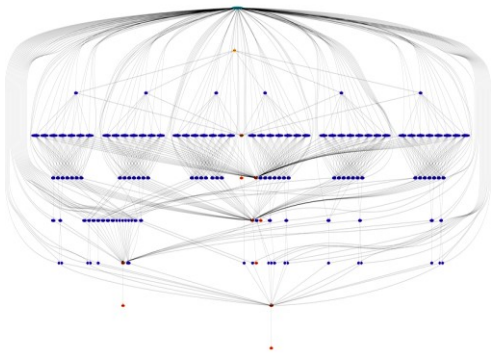

Becoming Data Scientists: Overcoming the Barriers

The goal of the class is to empower non-programmers
to communicate with computer scientists
so they can collaborate in real-world data science projects

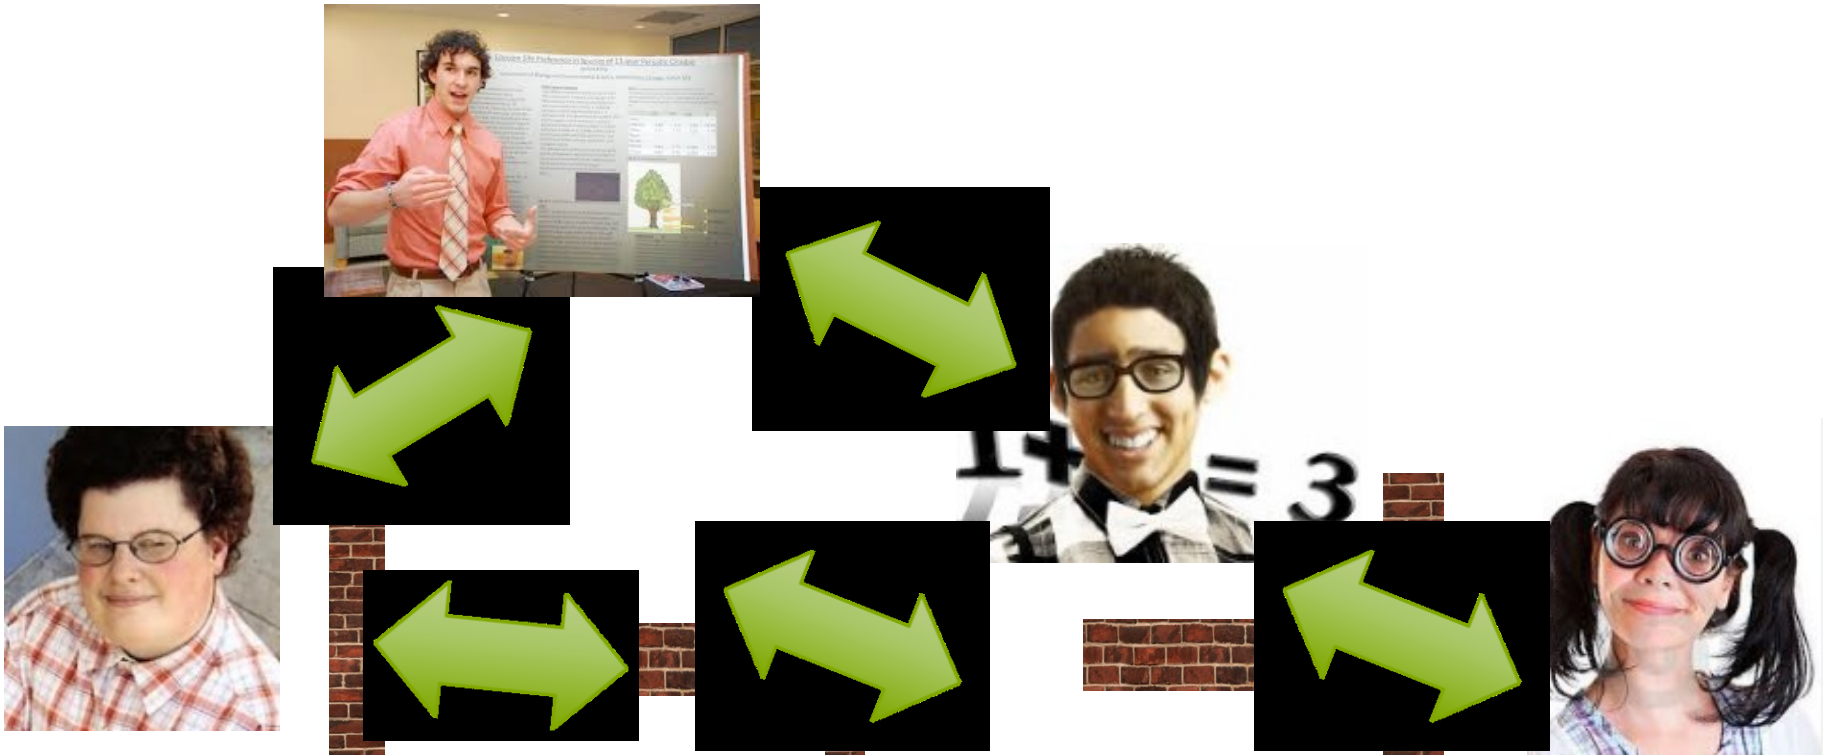

Students Learn to Channel Their Domain Expertise into Data Science

- USC course attended by graduate students in:
  - Political sciences, social sciences, education, biology, medicine, engineering
- Palpable trajectory:
  - Week 1: Sketch a data science project
    - Good goals, but impractical, nonsensical, unmanageable
  - Week 7: Revisit the data science project
    - Sensible approach, technical vocabulary
  - Week 12: Revisit again
    - Practical implementable approach

Wings workflow system

Defin questions

Collect/find data

Protégé ontology editor

collaborative projects

provenance

standards

Store data

ontologies

databases

Extract data

text, network, geospatial, timeseries

cleaning, imputation, reformatting

Pre-process data

parallel processing

Analyze data

classification, clustering, causality, simulation

workflows

Present results

visualization

repositories, licenses, DOIs

Publish data

Design Principles for the Course

Conceptual Learning

- **Computational thinking**: a new way to approach problems through computing
  - Simulation, data analysis, data mining
- **Data science**: a cross-disciplinary approach to solving data-rich problems
  - Machine learning, large-scale computing, semantic metadata

Practical Learning

- **Workflows**: a graphical programming environment that enables non-programmers to experiment with complex multi-step data analysis environments
- **Application domains**: exposure to past and ongoing projects where data science exposes multi-disciplinary challenges
  - Social networks, hydrology, proteomics, genomics, medicine, etc.

Practical Learning:
Workflows Make Data Science Accessible to Non-Programmers

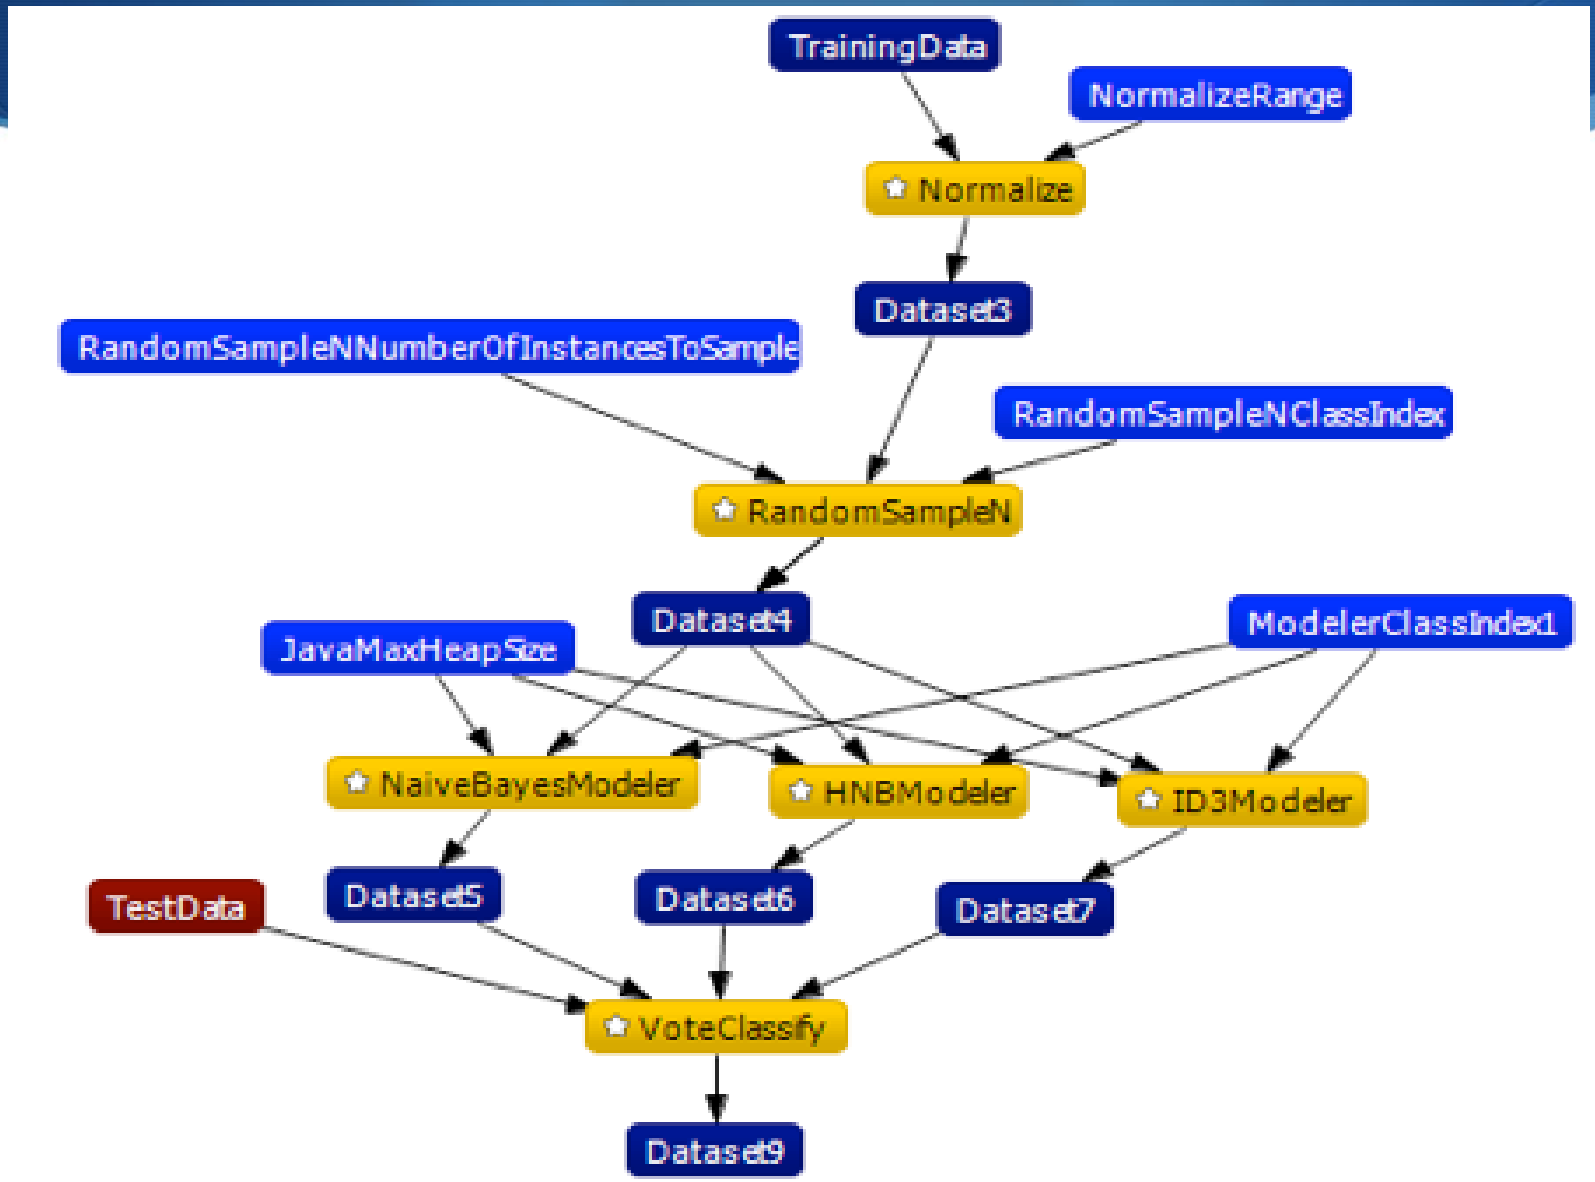

Practical Learning:
Workflows Enable Non-Programmers to Process Real-World Data

31	Scott Womack	1.073	00:47.25	00:46.29	3.600	Gold	Sedberry & Associates
33	Tucker VanOrmer	1.076	00:46.13	00:46.84	3.400	Gold	Red River Ski Team
60	Paul Lium	1.081	00:47.38	00:46.64	3.650	Gold	Jade Jets
51	Steve Dixon	1.086	00:46.58	00:47.56	3.617	Gold	Jade Jets
23	Mark Wilson	1.101	00:47.23	00:47.77	3.990	Gold	Jade Jets
28	Gary Vogel	1.114	00:47.93	00:48.06	3.967	Gold	C.O. Jones
45	Will Jackson	1.122	00:48.86	00:48.40	3.767	Gold	Jade Jets
61	Floater Floater	1.125	00:48.26	00:49.58	3.733	Gold	Red River Ski Team
64	Simon Mattsson	1.130	00:48.45	00:49.29	3.067	Gold	SNL SNo Problem
30	Guy Jackson	1.133	00:48.59	00:49.18	3.567	Gold	Guy Jackson &

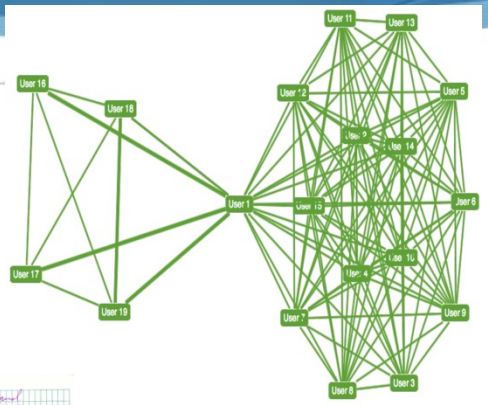

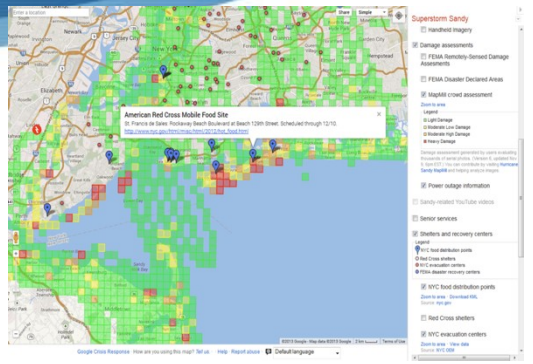

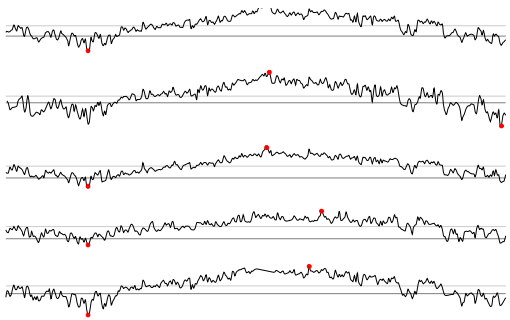

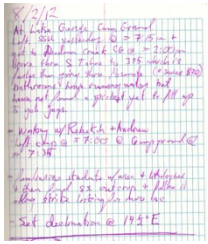

Boffa Surgical Pathology Report - Notepad

File Edit Format View Help

Cardenas Patient: SPRIGGS, KENNETH S :: 11:51

Page 2 of 2

D: Gallbladder

Gross Description:

A. The specimen is received fresh for intraoperative consultation. It consists of a resected right colon, portion of terminal ileum and portion of sigmoid colon. The right colon measures 50 cm in length, the terminal ileum segment measures 55 cm in length and the segment of sigmoid colon measures 7 cm in length. There are extensive serosal fibrous adhesions involving all three bowel segments; the ascending colon is distorted by side to side adhesions; there are side to side adhesions between two loops of terminal ileum and between a portion of terminal ileum and cecum/proximal right

colon; there are dense adhesions between the serosal surfaces of the distal

portion of the terminal ileum and the serosa of the sigmoid colon segment. There is a 0.3 cm diameter mucosal ulceration in the distal right colon, compatible with a biopsy site. The remainder of the right colon mucosa appears intact. The cecum is distorted and no appendix is identified. The terminal small bowel segment has a normal caliber at its proximal resection margin; the lumen is narrowed in the distal portion, but is not completely obstructed. The distal several cm of ileum show marked polypoid mucosal distortion with areas of irregular shallow mucosal ulceration; sections demonstrate fibrous thickening and scarring of the muscularis layers and serosa. Sectioning through the adhesed portion of terminal ileum and sigmoid colon demonstrate white fibrous scar tissue between the serosal surfaces of both bowel segments; however, a discreet fistulous tract is not identified with certainty. The following sections are submitted: 1-3 ascending colon

4-6 Terminal ileum and cecum

7 Proximal ileum, grossly unremarkable

8-10 Distal ileum with mucosal ulceration and polyps and fibrous scarring of muscularis

11&12 Distal ileum with side-to-side adhesions

13&14 Distal ileum and sigmoid with serosal adhesions

B. The specimen is received in formalin and consists of a 3.5 x 2.8 x 1.5 cm piece of yellow-gray adipose tissue partially covered on one side by a translucent gray serosa. Serial sections demonstrate focal yellow fat necrosis. Representative sections are submitted in one cassette.

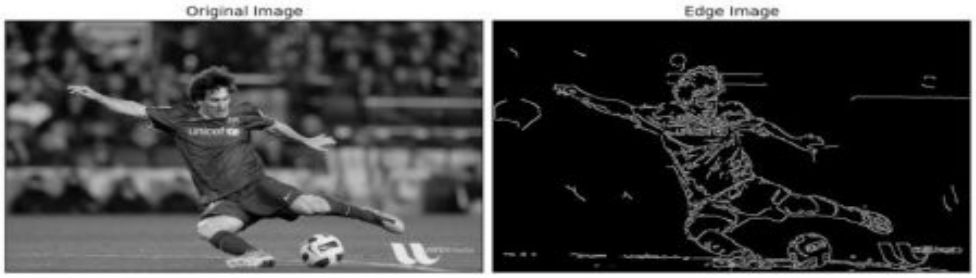

Experiment with Data Science in a Variety of Domains

image processing

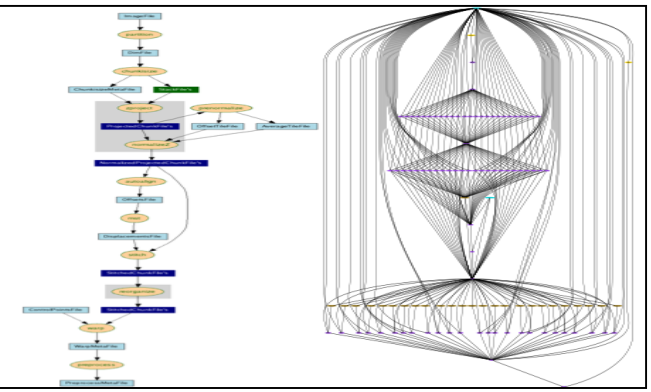

hydrology

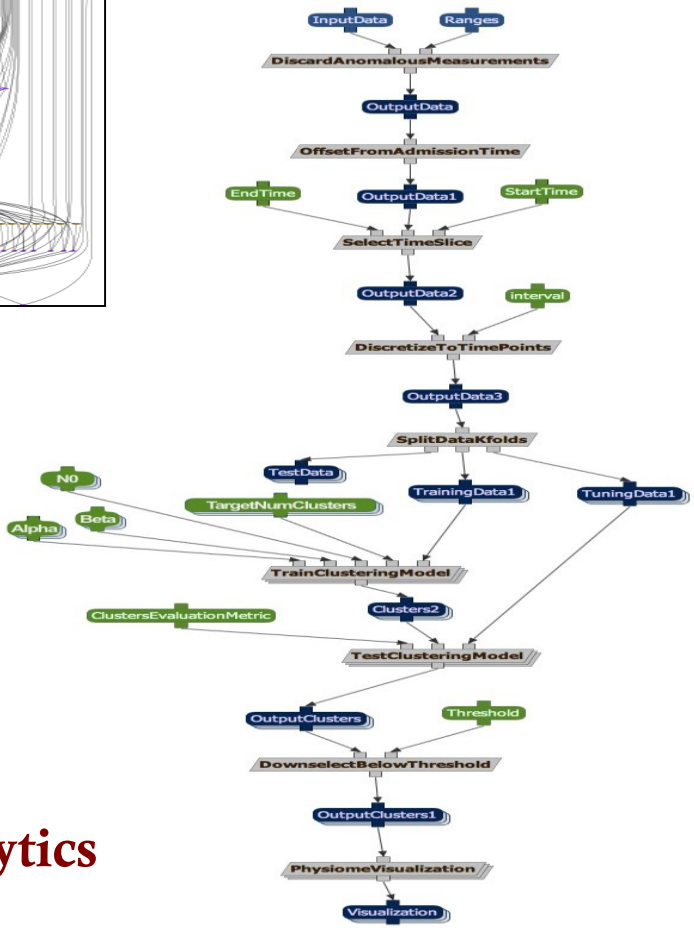

social network analysis

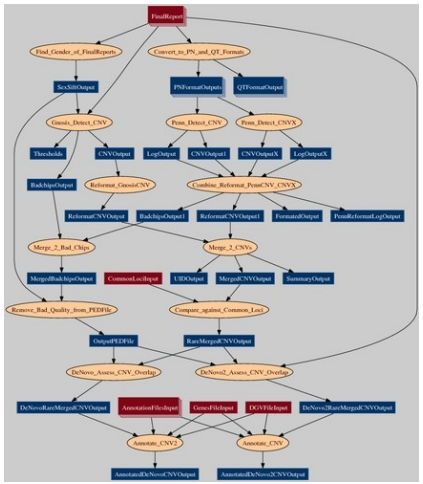

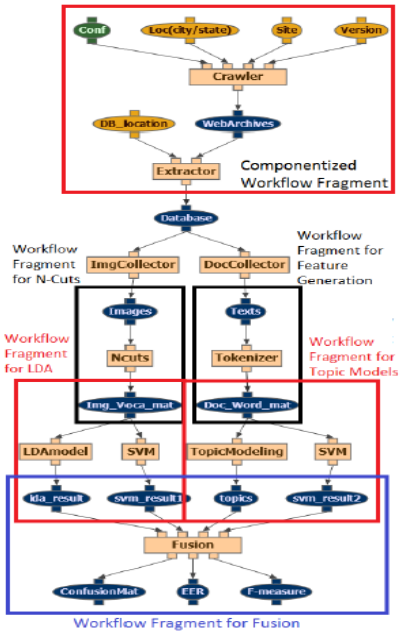

text analytics

biology

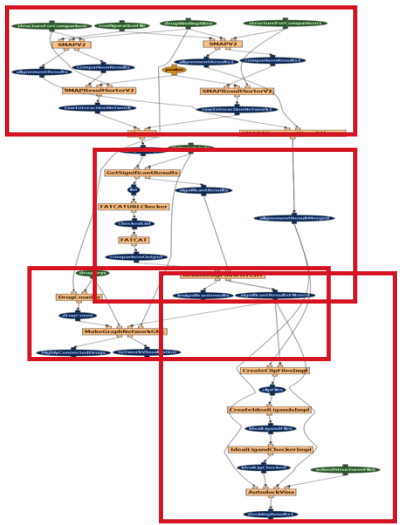

Section I: Introduction

1. Computational thinking and data science
2. Data
   - What is accessible data
   - Major types of data
   - Basic terminology
3. Data analysis software
   1. Algorithms vs code
   2. Programming languages
   3. Turing machines
4. Multi-step data analysis as workflows

http://blog.ginger.io/data-science-101/

II: Data Analysis

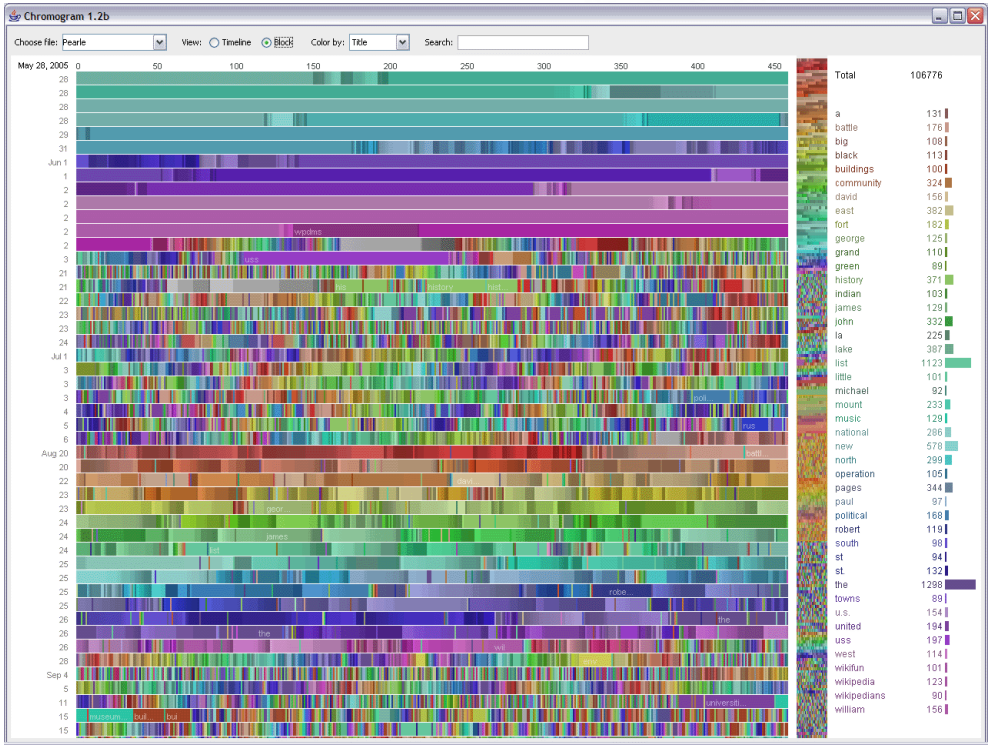

1. Data analysis tasks
   - Classification, clustering, pattern detection, causal discovery, simulation
2. Data pre-processing
3. Data visualization
4. Data lifecycle

https://es.wikipedia.org/wiki/Big_data

III: Data Analysis in Practice

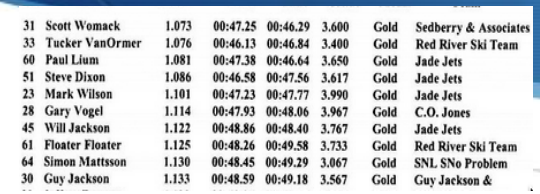

31	Scott Womack	1.073	00:47.25	00:46.29	3.600	Gold	Sedberry & Associates
33	Tucker VanOrmer	1.076	00:46.13	00:46.84	3.400	Gold	Red River Ski Team
60	Paul Lium	1.081	00:47.38	00:46.64	3.650	Gold	Jade Jets
51	Steve Dixon	1.086	00:46.58	00:47.56	3.617	Gold	Jade Jets
23	Mark Wilson	1.101	00:47.23	00:47.77	3.990	Gold	Jade Jets
28	Gary Vogel	1.114	00:47.93	00:48.06	3.967	Gold	C.O. Jones
45	Will Jackson	1.122	00:48.86	00:48.40	3.767	Gold	Jade Jets
61	Floater Floater	1.125	00:48.26	00:49.58	3.733	Gold	Red River Ski Team
64	Simon Mattsson	1.130	00:48.45	00:49.29	3.067	Gold	SNL SNo Problem
30	Guy Jackson	1.133	00:48.59	00:49.18	3.567	Gold	Guy Jackson &

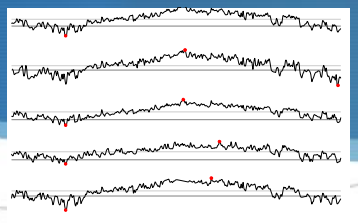

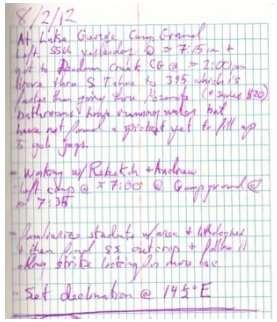

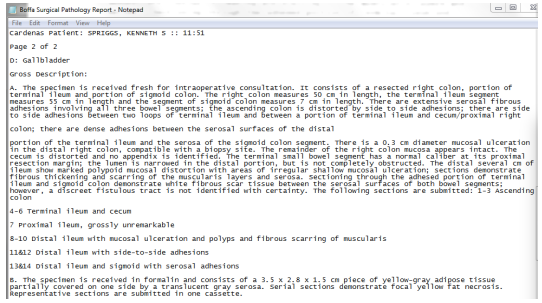

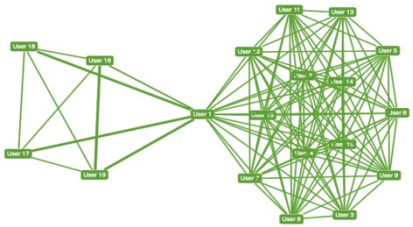

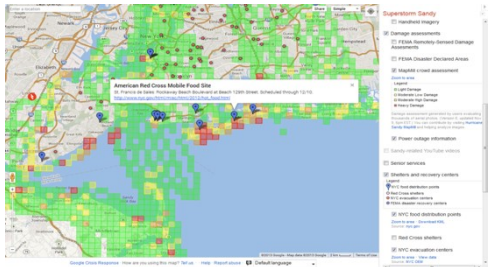

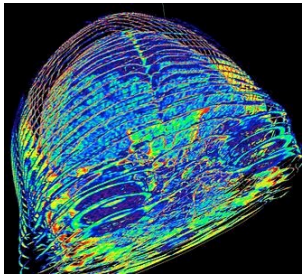

1. Analyzing different kinds of data
2. Parallel and distributed computing
   - Multi-core, clusters, grids, web services, …
   - Speedup, dependencies, critical paths, Amdahl's law, MapReduce

IV: Metadata

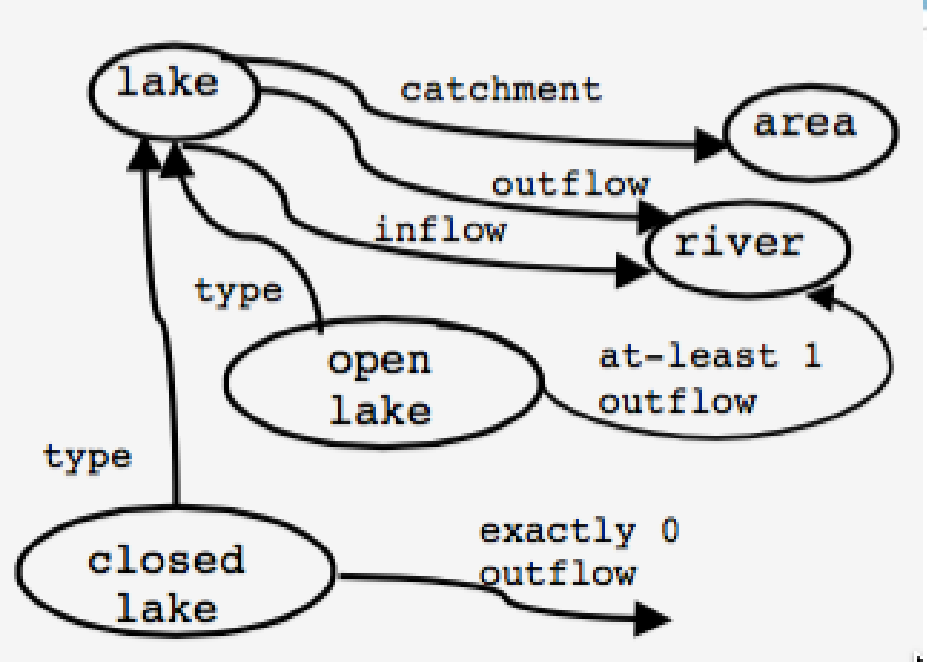

1. Semantic metadata
   - The DC standard
2. Ontologies
3. Provenance

V: Data Dissemination

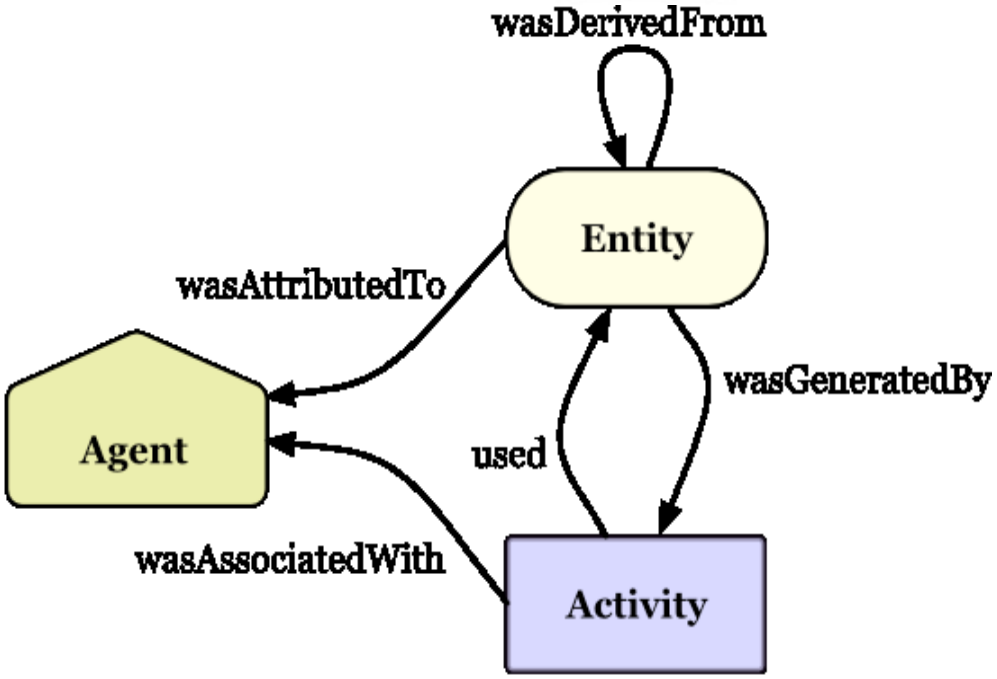

1. Data formats and standards
2. Provenance
3. Data stewardship
   - Data identifiers, data citation

VI: Advanced Topics

- Privacy and sensitive data
- Introduction to databases
- Crowdsourcing data collection
- Multi-disciplinary collaboration
- Project management

Course Design

Focus on AI Topics

- No statistics
- No databases
- No programming skills taught
- Yes: scalable algorithms

Constant Practice

- Class group activities followed by group reports
  - Learn to communicate and use technical terms
- Homeworks emphasize hypothesis formulation and testing

INF549 for USC Informatics Student Comments

- *"I went to a big data and big analytics keynote by a SAS VP. The presenter used a lot of jargon vocabulary that we have seen in class, and it was very interesting to see the application of many of the class topics. I understood the talk!"*
- *"I attended a Big Data student poster session in Engineering, and I could understand the presentations!"*

The prerequisites for the course were adequate.

1 Poor	0	0.00%
2 Below Average	0	0.00%
3 Average	1	6.67%
4 Above Average	5	33.33%
5 Excellent	9	60.00%
Total	15	

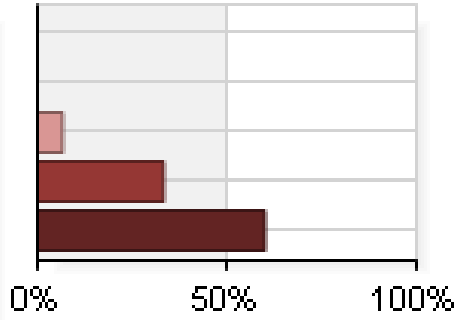

Encouraged students to participate in their learning (e.g., through discussion, projects, study groups and other appropriate activities).

1 Poor	0	0.00%
2 Below Average	0	0.00%
3 Average	0	0.00%
4 Above Average	1	6.25%
5 Excellent	15	93.75%
Total	16	

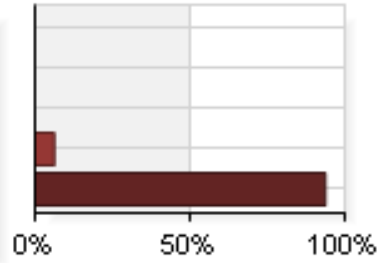

Ongoing

- Making materials available at datascience4all.org
  - Include videolectures
- Topical tutorials at science meetings (EarthCube, NOAA)
  - Eg, ontologies, machine learning
- North American Summer School in Data Science (with Caltech)
  - Already used for the 2016 RDA Summer School in Data Science
  - Already used in 2016 IS-GEO Summer School

ADDITIONAL SLIDES

Section I: Introduction to Basic Concepts

1. Computational thinking and data science
2. Data
   - Accessible data
     - APIs, license,…
   - Major types of data
     - Networks, text, time series, geospatial,…
   - Data terms
     - Metadata, silos, sensitive data, big data, …
3. Data analysis software
   - Algorithms vs code
     - Algorithm design
   - Programming languages
   - Encapsulation
   - Turing machines, Turing-complete languages
4. Multi-step data analysis as workflows
   - Components, dataflow, intermediate data
   - **Practicum: WINGS**

Section II: Data Analysis

1. Data analysis tasks
   - Classification
     - Decision trees
     - Alternative methods
     - Accuracy and other metrics
   - Pattern detection
     - Clustering
     - Temporal patterns
     - Network patterns
   - Causal discovery
     - Graphical models
     - Bayesian networks
   - Simulation

2. Data lifecycle
   - Data pre-processing
     - Data cleaning
   - Data wrangling
   - ETL
   - Collect, clean, analyze, visualize, deposit

3. Data visualization
   - Time series
   - Statistical
   - Maps, cartograms
   - Treemaps, heatmaps
   - Networks
   - Visual analytics

Section III: Data Analysis in Practice

1. Analyzing different kinds of data
   - Time series data
     - Ecology
     - Medicine
   - Text data
     - Web
     - Archives
   - Network data
     - Social media
     - Web
   - Multimedia data
     - Images
     - Videos
   - Geospatial data

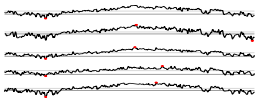

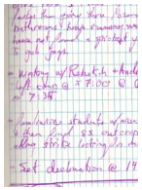

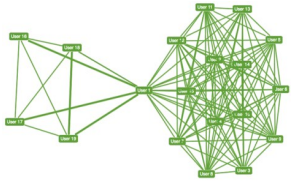

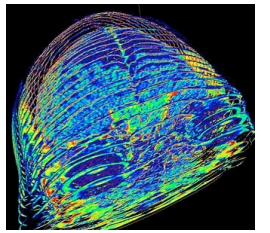

2. Parallel computing
   - Speedup estimates
   - Dependencies
   - Critical paths
   - Amdahl's law
   - MapReduce

3. Distributed computing
   - Multi-core computing, chips
   - Clusters
   - Grids
   - Web services
   - Cloud computing

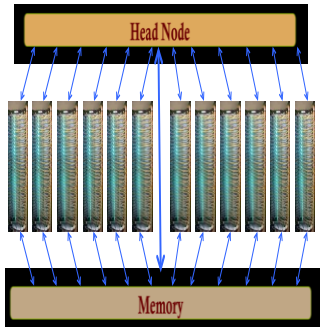

Section IV: Metadata

1. Semantic metadata
   - Attribution metadata
     - The Dublin Core standard
   - Summary metadata
   - Provenance metadata
   - Metadata in workflows
     - Metadata capture
     - Metadata propagation

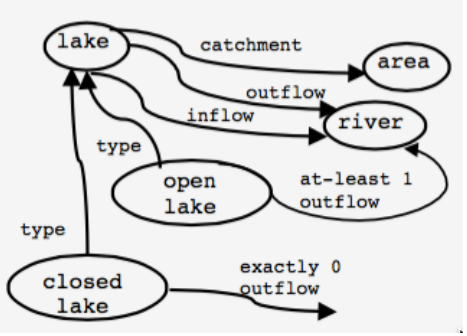

2. Ontologies
   - Taxonomies
   - Classes
   - Properties
   - Assertions
   - Definitions
   - Constraints and rules
   - Reasoning
   - **Practicum: PROTÉGÉ**

3. Provenance
   - Process provenance
   - Attribution provenance
   - Resource provenance
   - Provenance standards

Section V: Data Dissemination

1. Data formats
   - Data standards
   - Data repositories
   - Data services
   - Web data
2. Combining metadata and provenance
   - Metadata propagation
   - Automatic method validation
   - Automatic generation of metadata
   - Automatic provenance tracking
3. Data stewardship
   - Data sharing
   - Data identifiers
   - Licenses for data
   - Data citation and attribution
   - Software publication

VI: Advanced Topics

- Privacy and sensitive data
- Introduction to databases
- Crowdsourcing data collection
- Multi-disciplinary collaboration
- Project management